

# Assignment on predictive analysis

Chandrani Sengupta

2026-02-04

## Problem Set 2: Linear Regression

### Question 1 : Problem to demonstrate that the population regression line is fixed, but least square regression line varies

Suppose the population regression line is given by  $Y = 2 + 3x$ , while the data comes from the model  $y = 2 + 3x + \varepsilon$ . Step 1: For  $x$  in the range [5,10] graph the population regression line. Step 2: Generate  $x_i(i = 1, 2, \dots, n)$  from Uniform(5, 10) and  $\varepsilon_i(i = 1, 2, \dots, n)$  from  $N(0, 4.2)$ . Hence, compute  $y_1, y_2, \dots, y_n$ . Step 3: On the basis of the data  $(x_i, y_i)(i = 1, 2, \dots, n)$  generated in Step 2,

report the least squares regression line. Step 4: Repeat steps 2-3 five times. Graph the 5 least squares regression lines over the population regression line obtained in Step 1.

Interpret the findings. Take  $n = 50$ . Set the seed as seed=123.

```
#step 1
x=seq(5,10,length.out=200)
x

## [1] 5.000000 5.025126 5.050251 5.075377 5.100503 5.125628
5.150754
## [8] 5.175879 5.201005 5.226131 5.251256 5.276382 5.301508
5.326633
## [15] 5.351759 5.376884 5.402010 5.427136 5.452261 5.477387
5.502513
## [22] 5.527638 5.552764 5.577889 5.603015 5.628141 5.653266
5.678392
## [29] 5.703518 5.728643 5.753769 5.778894 5.804020 5.829146
5.854271
## [36] 5.879397 5.904523 5.929648 5.954774 5.979899 6.005025
6.030151
## [43] 6.055276 6.080402 6.105528 6.130653 6.155779 6.180905
6.206030
## [50] 6.231156 6.256281 6.281407 6.306533 6.331658 6.356784
6.381910
## [57] 6.407035 6.432161 6.457286 6.482412 6.507538 6.532663
6.557789
## [64] 6.582915 6.608040 6.633166 6.658291 6.683417 6.708543
6.733668
## [71] 6.758794 6.783920 6.809045 6.834171 6.859296 6.884422
6.909548
```

```

## [78] 6.934673 6.959799 6.984925 7.010050 7.035176 7.060302
7.085427
## [85] 7.110553 7.135678 7.160804 7.185930 7.211055 7.236181
7.261307
## [92] 7.286432 7.311558 7.336683 7.361809 7.386935 7.412060
7.437186
## [99] 7.462312 7.487437 7.512563 7.537688 7.562814 7.587940
7.613065
## [106] 7.638191 7.663317 7.688442 7.713568 7.738693 7.763819
7.788945
## [113] 7.814070 7.839196 7.864322 7.889447 7.914573 7.939698
7.964824
## [120] 7.989950 8.015075 8.040201 8.065327 8.090452 8.115578
8.140704
## [127] 8.165829 8.190955 8.216080 8.241206 8.266332 8.291457
8.316583
## [134] 8.341709 8.366834 8.391960 8.417085 8.442211 8.467337
8.492462
## [141] 8.517588 8.542714 8.567839 8.592965 8.618090 8.643216
8.668342
## [148] 8.693467 8.718593 8.743719 8.768844 8.793970 8.819095
8.844221
## [155] 8.869347 8.894472 8.919598 8.944724 8.969849 8.994975
9.020101
## [162] 9.045226 9.070352 9.095477 9.120603 9.145729 9.170854
9.195980
## [169] 9.221106 9.246231 9.271357 9.296482 9.321608 9.346734
9.371859
## [176] 9.396985 9.422111 9.447236 9.472362 9.497487 9.522613
9.547739
## [183] 9.572864 9.597990 9.623116 9.648241 9.673367 9.698492
9.723618
## [190] 9.748744 9.773869 9.798995 9.824121 9.849246 9.874372
9.899497
## [197] 9.924623 9.949749 9.974874 10.000000

```

$y=2+3*x$

y

```

## [1] 17.00000 17.07538 17.15075 17.22613 17.30151 17.37688 17.45226
17.52764
## [9] 17.60302 17.67839 17.75377 17.82915 17.90452 17.97990 18.05528
18.13065
## [17] 18.20603 18.28141 18.35678 18.43216 18.50754 18.58291 18.65829
18.73367
## [25] 18.80905 18.88442 18.95980 19.03518 19.11055 19.18593 19.26131
19.33668
## [33] 19.41206 19.48744 19.56281 19.63819 19.71357 19.78894 19.86432
19.93970
## [41] 20.01508 20.09045 20.16583 20.24121 20.31658 20.39196 20.46734

```

```

20.54271
## [49] 20.61809 20.69347 20.76884 20.84422 20.91960 20.99497 21.07035
21.14573
## [57] 21.22111 21.29648 21.37186 21.44724 21.52261 21.59799 21.67337
21.74874
## [65] 21.82412 21.89950 21.97487 22.05025 22.12563 22.20101 22.27638
22.35176
## [73] 22.42714 22.50251 22.57789 22.65327 22.72864 22.80402 22.87940
22.95477
## [81] 23.03015 23.10553 23.18090 23.25628 23.33166 23.40704 23.48241
23.55779
## [89] 23.63317 23.70854 23.78392 23.85930 23.93467 24.01005 24.08543
24.16080
## [97] 24.23618 24.31156 24.38693 24.46231 24.53769 24.61307 24.68844
24.76382
## [105] 24.83920 24.91457 24.98995 25.06533 25.14070 25.21608 25.29146
25.36683
## [113] 25.44221 25.51759 25.59296 25.66834 25.74372 25.81910 25.89447
25.96985
## [121] 26.04523 26.12060 26.19598 26.27136 26.34673 26.42211 26.49749
26.57286
## [129] 26.64824 26.72362 26.79899 26.87437 26.94975 27.02513 27.10050
27.17588
## [137] 27.25126 27.32663 27.40201 27.47739 27.55276 27.62814 27.70352
27.77889
## [145] 27.85427 27.92965 28.00503 28.08040 28.15578 28.23116 28.30653
28.38191
## [153] 28.45729 28.53266 28.60804 28.68342 28.75879 28.83417 28.90955
28.98492
## [161] 29.06030 29.13568 29.21106 29.28643 29.36181 29.43719 29.51256
29.58794
## [169] 29.66332 29.73869 29.81407 29.88945 29.96482 30.04020 30.11558
30.19095
## [177] 30.26633 30.34171 30.41709 30.49246 30.56784 30.64322 30.71859
30.79397
## [185] 30.86935 30.94472 31.02010 31.09548 31.17085 31.24623 31.32161
31.39698
## [193] 31.47236 31.54774 31.62312 31.69849 31.77387 31.84925 31.92462
32.00000

plot(x,y,type='l')
#step2
n=50
x1=runif(50,5,10)
x1

## [1] 7.969716 9.802671 8.558071 5.444186 7.459862 9.687883 8.216514
6.351221
## [9] 5.515075 8.908464 5.989649 9.376409 5.204452 6.854462 9.635217
7.681841

```

```

## [17] 9.250357 8.926526 6.652887 9.659516 8.623565 8.794105 5.243854
6.672255
## [25] 8.730931 8.354203 7.598708 5.726385 8.590139 7.360572 6.846349
7.262759
## [33] 9.451626 8.625439 6.165179 6.734055 6.619490 7.273549 5.021913
7.152649
## [41] 7.409838 7.627102 7.708101 5.634432 6.094051 9.386364 7.430531
8.982201
## [49] 5.325925 8.343973

e=rnorm(50,0,4)
e

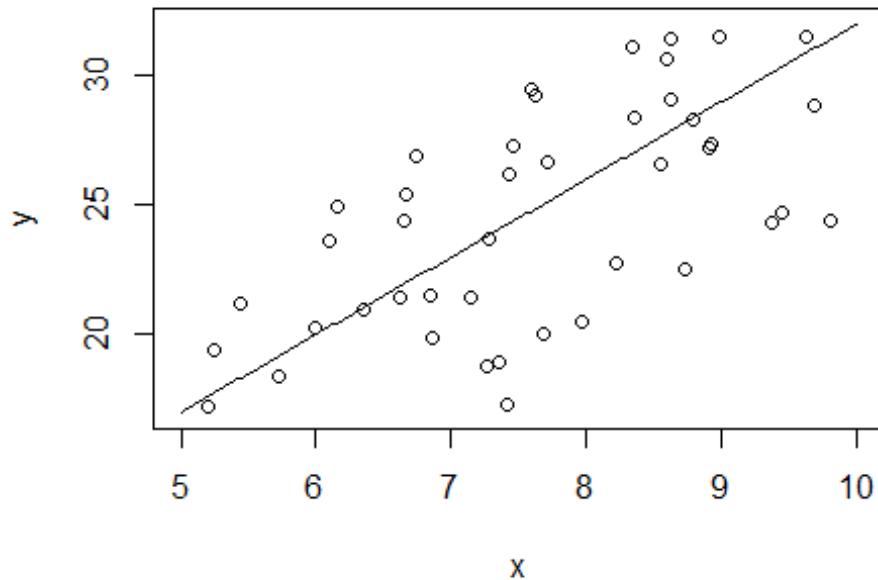
## [1] -5.39343803 -7.01921557 -1.07544142 2.85633781 2.96057313 -
2.16072507
## [7] -3.85190811 -0.07078827 -9.55620984 -1.53661675 0.27875533 -
5.81085605
## [13] -0.41751804 -2.65413106 0.65266939 -5.00654445 5.63594574 -
1.39750081
## [19] 2.41738215 9.10290410 3.56500631 -0.04186448 1.70791937
3.43054471
## [25] -5.65629232 1.33070910 4.68242828 -0.78183322 2.93033457 -
5.11923843
## [31] -1.03321432 -4.98856727 -5.60022368 1.26181681 4.42633547
4.74308947
## [37] -0.43367473 -0.14537730 -1.84958336 -2.04136103 -6.91170270
4.39029483
## [43] 1.55226501 -3.08342295 3.35827158 6.71350113 1.88409998
2.60088200
## [49] -5.26477367 4.11314241

y1=2+3*x1+e
y1

## [1] 20.515709 24.388799 26.598771 21.188896 27.340159 28.902924 22.797633
## [8] 20.982874 8.989014 27.188775 20.247701 24.318371 17.195837 19.909255
## [15] 31.558320 20.038980 35.387017 27.382078 24.376044 40.081452 31.435701
## [22] 28.340451 19.439482 25.447308 22.536502 28.393318 29.478552 18.397322
## [29] 30.700752 18.962478 21.505831 18.799710 24.754654 29.138134 24.921873
## [36] 26.945256 21.424797 23.675270 15.216157 21.416587 17.317810 29.271602
## [43] 26.676568 15.819872 23.640424 36.872594 26.175694 31.547484 12.713001
## [50] 31.145061

points(x1,y1)

```



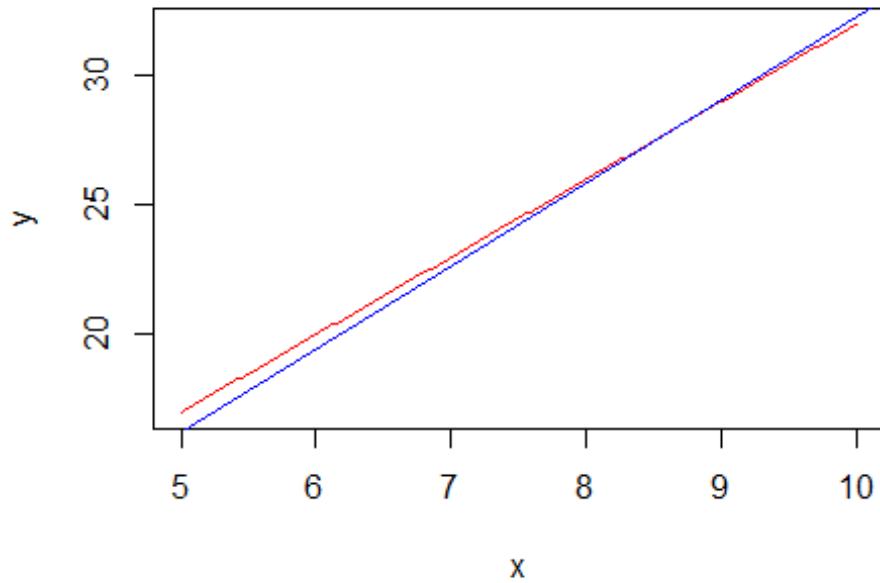
```
#step3
lin.reg=lm(y1~x1)
lin.reg#Lm object

##
## Call:
## lm(formula = y1 ~ x1)
##
## Coefficients:
## (Intercept)      x1
##       0.1032      3.2184

coef(lin.reg)

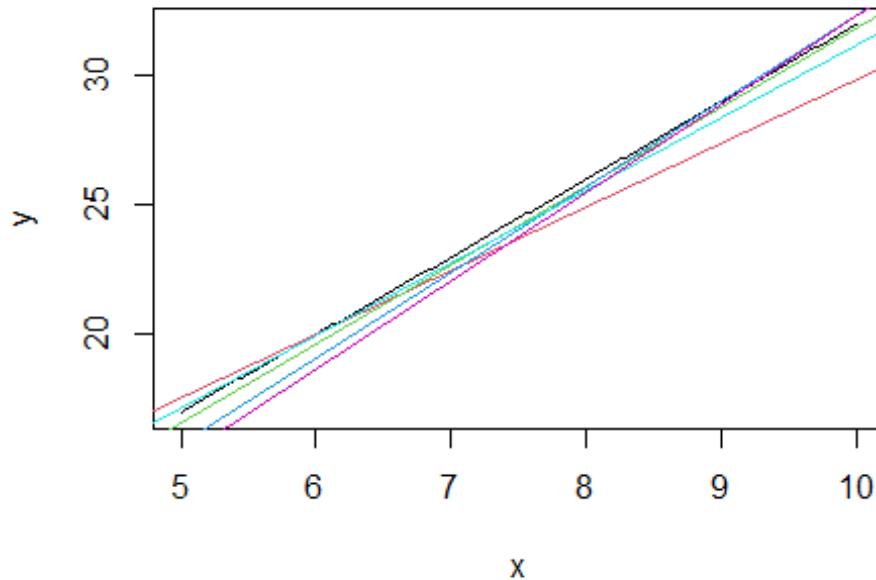
## (Intercept)      x1
##   0.1031788    3.2183820

plot(x,y,type='l',col="red")#population regression line
abline(lin.reg,col="blue")#sample regression line
```



```
#step4
plot(x,y,type='l')
n=50
beta.hat=c()
for(i in 1:5){

  x2=rnorm(n,5,10)
  eps=rnorm(n,0,4)
  y2=2+3*x2+eps
  lin.reg=lm(y2~x2)
  beta.hat=rbind(beta.hat,coef(lin.reg))
  abline(lin.reg,col=i+1)
}
```



```
beta.hat=as.data.frame(beta.hat,colnames=c("beta_0_hat","beta_hat"))
head(beta.hat)

##      (Intercept)      x2
## 1  5.2033775 2.466752
## 2  1.3701508 3.046174
## 3 -0.8241796 3.314432
## 4  3.0977081 2.812735
## 5 -1.7479371 3.403991
```

## Question 2 : Problem to demonstrate that $\hat{\beta}_0$ and $\hat{\beta}$ minimises RSS

Step 1: Generate  $x_i$  from Uniform(5, 10) and mean centre the values. Generate  $\epsilon_i$  from  $N(0, 1)$ . Calculate  $y_i = 2 + 3x_i + \epsilon_i$

,  $i = 1, 2, \dots, n$ . Take  $n=50$  and

`seed=123.` Step 2: Now imagine that you only have the data on  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ , without knowing the mechanism that was used to generate the data in step 1. Assuming a linear regression of the type  $y_i = \beta_0 + \beta x_i + \epsilon_i$

, and based on these

`data (xi , yi), i = 1, 2, ..., n,` obtain the least squares estimates of  $\beta_0$  and  $\beta$ . Step 3: Take a large number of grid values of  $(\beta_0, \beta)$  that also include the least squares estimates

obtained from step 2. Compute the RSS for each parametric choice of  $(\beta_0, \beta)$ , where  $\text{RSS} = (y_1 - \beta_0 - \beta x_1)^2 + (y_2 - \beta_0 - \beta x_2)^2 + \dots + (y_n - \beta_0 - \beta x_n)^2$ . Find out for which combination of  $(\beta_0, \beta)$ , RSS is minimum.

```
2 + (y2 - β0 - βx2)
```

```
2 + ....(yn -
```

$\beta_0 - \beta x_n)^2$ . Find out for which combination of  $(\beta_0, \beta)$ , RSS is minimum.

```
#step1
```

```
x=runif(50,5,10)
x

## [1] 6.184547 7.960892 5.030061 9.148818 5.543152 7.283953 6.409737
7.404770
## [9] 7.092128 5.455692 7.403749 9.214539 6.065383 5.783260 9.529368
6.295444
## [17] 8.863508 8.554613 6.830856 6.494904 5.041425 6.720073 8.862320
5.999880
## [25] 5.281977 6.293306 6.956423 7.254962 7.309258 6.291363 8.837801
6.221546
## [33] 9.814047 5.456501 7.103373 9.967378 5.713610 8.135989 6.959268
8.956002
## [41] 5.514825 9.899958 7.269552 9.570969 9.998060 9.763062 7.861456
6.379462
## [49] 6.006276 6.140459

eps=rnorm(50,0,1)
eps

## [1] -0.985702853 -0.766779668 -0.373632231 -0.960611817 0.435824622
## [6] -0.545872976 -1.681042607 0.077959405 -0.718442729 -2.106568039
## [11] -0.655608394 0.032740067 -1.017633524 -1.878365874 -0.006234435
## [16] -1.022627806 -0.789217131 -0.737715214 -0.005511238 1.267501570
## [21] 0.494603528 -0.136741054 2.518204651 -0.028904409 0.114965346
## [26] -0.715999897 -2.126166656 0.829429548 0.417751703 -0.214898704
## [31] 0.309045497 0.561123822 0.019602620 -0.072032936 -1.677541983
## [36] 0.652685945 0.555004334 0.077208314 -1.509826736 -0.369157003
## [41] -1.983371022 0.329924145 1.270275459 -1.434768413 -0.425525920
## [46] 1.100934387 -0.242491378 1.152271054 1.150912154 0.207809105
```

```
#step2
```

```
x_mean=mean(x)
x_mean_centered=x-x_mean
x_mean_centered

## [1] -1.098052103 0.678293105 -2.252537773 1.866218815 -1.739447426
## [6] 0.001354336 -0.872862049 0.122171194 -0.190471496 -1.826907045
## [11] 0.121149635 1.931939975 -1.217215893 -1.499339117 2.246768529
## [16] -0.987154808 1.580909426 1.272014087 -0.451743384 -0.787694831
```

```

## [21] -2.241174189 -0.562526459  1.579720858 -1.282718660 -2.000622164
## [26] -0.989292902 -0.326176230 -0.027637217  0.026659231 -0.991236189
## [31]  1.555201745 -1.061053181  2.531448137 -1.826098537 -0.179226244
## [36]  2.684779043 -1.568988974  0.853389481 -0.323331365  1.673402443
## [41] -1.767773729  2.617358581 -0.013047532  2.288369926  2.715461142
## [46]  2.480463183  0.578857024 -0.903136649 -1.276323283 -1.142140469

y=2+3*x_mean_centered+eps
y

## [1] -2.2798592  3.2680996 -5.1312455  6.6380446 -2.7825177  1.4581900
## [7] -2.2996288  2.4444730  0.7101428 -5.5872892  1.7078405  7.8285600
## [13] -2.6692812 -4.3763832  8.7340712 -1.9840922  5.9535111  5.0783270
## [19]  0.6392586  0.9044171 -4.2289190  0.1756796  9.2573672 -1.8770604
## [25] -3.8869011 -1.6838786 -1.1046953  2.7465179  2.4977294 -1.1886073
## [31]  6.9746507 -0.6220357  9.6139470 -3.5503285 -0.2152207 10.7070231
## [37] -2.1519626  4.6373768 -0.4798208  6.6510503 -5.2866922 10.1819999
## [43]  3.2311329  7.4303414  9.7208575 10.5423239  3.4940797  0.4428611
## [49] -0.6780577 -1.2186123

lin.reg=lm(y~x)
lin.reg

##
## Call:
## lm(formula = y ~ x)
##
## Coefficients:
## (Intercept)          x
##      -20.833        3.103

beta0_hat=coef(lin.reg) [1]
beta_hat=coef(lin.reg)[2]

beta0_hat

## (Intercept)
## -20.83345

beta_hat

##          x
## 3.103451

beta0_grid= seq(beta0_hat-2, beta0_hat + 2, length.out = 100)
beta_grid = seq(beta_hat-2, beta_hat + 2, length.out = 100)

RSS <- matrix(NA, nrow = length(beta0_grid), ncol = length(beta_grid))

# Compute RSS for each (beta0, beta)
for (i in 1:length(beta0_grid)) {

```

```

for (j in 1:length(beta_grid)) {
  RSS[i, j] <- sum((y1 - beta0_grid[i] - beta_grid[j] * x1)^2)
}
which(RSS == min(RSS), arr.ind=TRUE)

##      row col
## [1,] 100 100

# Find minimum RSS
min_RSS <- min(RSS)
index <- which(RSS == min_RSS, arr.ind = TRUE)
index

##      row col
## [1,] 100 100

```

##Question 3 : Problem to demonstrate that least square estimators are unbiased

Step 1: Generate  $x_i$ ( $i = 1, 2, \dots, n$ ) from Uniform(0, 1),  $\varepsilon_i$ ( $i = 1, 2, \dots, n$ ) from  $N(0, 1)$  and hence generate  $y$  using  $y_i = \beta_0 + \beta x_i + \varepsilon_i$

. (Take  $\beta_0 = 2$ ,  $\beta = 3$ ).

Step 2: On the basis of the data ( $x_i$

,  $y_i$ )( $i = 1, 2, \dots, n$ ) generated in Step 1,

obtain the least square estimates of  $\beta_0$  and  $\beta$ . Repeat Steps 1-2,  $R = 1000$  times. In each simulation obtain  $\hat{\beta}$

0 and  $\hat{\beta}$ . Finally, the least-square estimates will be given by the average of these estimated values. Compare these with the true  $\beta_0$  and  $\beta$  and comment. Take  $n = 50$  and seed=123.

```

x=runif(50,0,1)
x

## [1] 0.91954030 0.43609056 0.55652145 0.63723548 0.49624580 0.78449462
## [7] 0.44903243 0.66848769 0.29980540 0.88313706 0.03624314 0.49694234
## [13] 0.85768868 0.08043533 0.13068987 0.94042848 0.08242313 0.95761017
## [19] 0.37303280 0.43664158 0.59222158 0.55077948 0.37940802 0.67552168
## [25] 0.96426716 0.41476347 0.59516219 0.64023404 0.55318780 0.19619461
## [31] 0.29200771 0.41083248 0.79893475 0.41551168 0.77732370 0.50439703
## [37] 0.75579330 0.18213116 0.19591294 0.47069061 0.53980904 0.38129912
## [43] 0.07102295 0.41714622 0.05595624 0.19683692 0.71674888 0.89386611
## [49] 0.78749603 0.34370396

eps=rnorm(50,0,1)
eps

```

```

## [1] -0.069234486 -0.469155673 -1.281806623 1.550826371 1.306605699
## [6] 1.617492140 -1.011527050 0.736505744 -0.724979760 -0.419388147
## [11] 2.482719838 -0.502607181 1.021796486 0.468301839 -0.951487135
## [16] -1.024428388 -0.272726842 -0.936856090 0.683643418 -0.531612523
## [21] 1.296531873 1.522019019 -0.490273153 1.939867799 -0.372639001
## [26] 0.678443614 0.001943419 0.454984054 0.374740448 -1.024112288
## [31] 0.490215698 1.542885311 1.487311429 1.578516412 0.392557275
## [36] 1.368257224 0.707285334 0.596229524 0.799482749 0.545883459
## [41] -0.443146164 -0.704946297 0.516925883 -1.154832531 0.099309907
## [46] -1.151124845 2.221067283 0.637563328 -1.269752883 -0.929544778

y=2+3*x+eps
y

## [1] 4.689386 2.839116 2.387758 5.462533 4.795343 5.970976 2.335570
4.741969
## [9] 2.174436 4.230023 4.591449 2.988220 5.594863 2.709608 1.440582
3.796857
## [17] 1.974543 3.935974 3.802742 2.778312 5.073197 5.174357 2.647951
5.966433
## [25] 4.520162 3.922734 3.787430 4.375686 4.034304 1.564472 3.366239
4.775383
## [33] 5.884116 4.825051 4.724528 4.881448 4.974665 3.142623 3.387222
3.957955
## [41] 3.176281 2.438951 2.729995 2.096606 2.267179 1.439386 6.371314
5.319162
## [49] 3.092735 2.101567

lin.reg=lm(y~x)
lin.reg

##
## Call:
## lm(formula = y ~ x)
##
## Coefficients:
## (Intercept)           x
##       2.170            3.194

beta.hat=c()
for(i in 1:1000)
{
  x=runif(50,0,1)
  eps=rnorm(50,0,1)
  y=2+3*x+eps
  lin.reg=lm(y~x)
  beta.hat=rbind(beta.hat,coef(lin.reg))
}
head(beta.hat)

```

```

##      (Intercept)      x
## [1,] 1.516681 3.443819
## [2,] 1.931161 3.107286
## [3,] 2.008123 2.982416
## [4,] 2.423518 2.835003
## [5,] 1.670347 3.660578
## [6,] 2.094118 2.563321

avg_beta0=mean(beta.hat[,1])
avg_beta0

## [1] 1.988911

avg_beta=mean(beta.hat[,2])
avg_beta

## [1] 3.001523

```

## Question 4 : Comparing several simple linear regressions

Attach “Boston” data from MASS library in R. Select median value of owner- occupied homes, as the response and per capita crime rate, nitrogen oxides

concentration, proportion of blacks and percentage of lower status of the population as predictors.

- (a) Selecting the predictors one by one, run four separate linear regressions to the data. Present the output in a single table.
- (b) Which model gives the best fit?
- (c) Compare the coefficients of the predictors from each model and comment on the usefulness of the predictors.

```

library(MASS)

## Warning: package 'MASS' was built under R version 4.5.2

x1=Boston
y=Boston$medv#response
x=data.frame(x1$crim,x1$nox,x1$black,x1$lstat)#predictor
y

## [1] 24.0 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15.0 18.9 21.7
20.4 18.2
## [16] 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8
18.4 21.0
## [31] 12.7 14.5 13.2 13.1 13.5 18.9 20.0 21.0 24.7 30.8 34.9 26.6 25.3
24.7 21.2
## [46] 19.3 20.0 16.6 14.4 19.4 19.7 20.5 25.0 23.4 18.9 35.4 24.7 31.6
23.3 19.6
## [61] 18.7 16.0 22.2 25.0 33.0 23.5 19.4 22.0 17.4 20.9 24.2 21.7 22.8
23.4 24.1

```

```

## [76] 21.4 20.0 20.8 21.2 20.3 28.0 23.9 24.8 22.9 23.9 26.6 22.5 22.2
23.6 28.7
## [91] 22.6 22.0 22.9 25.0 20.6 28.4 21.4 38.7 43.8 33.2 27.5 26.5 18.6
19.3 20.1
## [106] 19.5 19.5 20.4 19.8 19.4 21.7 22.8 18.8 18.7 18.5 18.3 21.2 19.2
20.4 19.3
## [121] 22.0 20.3 20.5 17.3 18.8 21.4 15.7 16.2 18.0 14.3 19.2 19.6 23.0
18.4 15.6
## [136] 18.1 17.4 17.1 13.3 17.8 14.0 14.4 13.4 15.6 11.8 13.8 15.6 14.6
17.8 15.4
## [151] 21.5 19.6 15.3 19.4 17.0 15.6 13.1 41.3 24.3 23.3 27.0 50.0 50.0
50.0 22.7
## [166] 25.0 50.0 23.8 23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6
29.9 37.2
## [181] 39.8 36.2 37.9 32.5 26.4 29.6 50.0 32.0 29.8 34.9 37.0 30.5 36.4
31.1 29.1
## [196] 50.0 33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50.0 22.6 24.4 22.5
24.4 20.0
## [211] 21.7 19.3 22.4 28.1 23.7 25.0 23.3 28.7 21.5 23.0 26.7 21.7 27.5
30.1 44.8
## [226] 50.0 37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29.0 24.0 25.1 31.5
23.7 23.3
## [241] 22.0 20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8 29.6
42.8 21.9
## [256] 20.9 44.0 50.0 36.0 30.1 33.8 43.1 48.8 31.0 36.5 22.8 30.7 50.0
43.5 20.7
## [271] 21.1 25.2 24.4 35.2 32.4 32.0 33.2 33.1 29.1 35.1 45.4 35.4 46.0
50.0 32.2
## [286] 22.0 20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9 21.7 28.6 27.1 20.3
22.5 29.0
## [301] 24.8 22.0 26.4 33.1 36.1 28.4 33.4 28.2 22.8 20.3 16.1 22.1 19.4
21.6 23.8
## [316] 16.2 17.8 19.8 23.1 21.0 23.8 23.1 20.4 18.5 25.0 24.6 23.0 22.2
19.3 22.6
## [331] 19.8 17.1 19.4 22.2 20.7 21.1 19.5 18.5 20.6 19.0 18.7 32.7 16.5
23.9 31.2
## [346] 17.5 17.2 23.1 24.5 26.6 22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7
22.7 22.6
## [361] 25.0 19.9 20.8 16.8 21.9 27.5 21.9 23.1 50.0 50.0 50.0 50.0 50.0
13.8 13.8
## [376] 15.0 13.9 13.3 13.1 10.2 10.4 10.9 11.3 12.3 8.8 7.2 10.5 7.4
10.2 11.5
## [391] 15.1 23.2 9.7 13.8 12.7 13.1 12.5 8.5 5.0 6.3 5.6 7.2 12.1
8.3 8.5
## [406] 5.0 11.9 27.9 17.2 27.5 15.0 17.2 17.9 16.3 7.0 7.2 7.5 10.4
8.8 8.4
## [421] 16.7 14.2 20.8 13.4 11.7 8.3 10.2 10.9 11.0 9.5 14.5 14.1 16.1
14.3 11.7
## [436] 13.4 9.6 8.7 8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
14.1 13.0

```

```

## [451] 13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20.0 16.4 17.7 19.5
20.2 21.4
## [466] 19.9 19.0 19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3 16.7 12.0
14.6 21.4
## [481] 23.0 23.7 25.0 21.8 20.6 21.2 19.1 20.6 15.2 7.0 8.1 13.6 20.1
21.8 24.5
## [496] 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9 22.0 11.9

model1=lm(y~x1$crim)
model1

##
## Call:
## lm(formula = y ~ x1$crim)
##
## Coefficients:
## (Intercept)      x1$crim
##        24.0331     -0.4152

summary(model1)

##
## Call:
## lm(formula = y ~ x1$crim)
##
## Residuals:
##    Min     1Q   Median     3Q    Max
## -16.957 -5.449 -2.007  2.512 29.800
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 24.03311   0.40914  58.74 <2e-16 ***
## x1$crim     -0.41519   0.04389  -9.46 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

## 
## Residual standard error: 8.484 on 504 degrees of freedom
## Multiple R-squared:  0.1508, Adjusted R-squared:  0.1491
## F-statistic: 89.49 on 1 and 504 DF,  p-value: < 2.2e-16

model2=lm(y~x1$nox)
model2

##
## Call:
## lm(formula = y ~ x1$nox)
##
## Coefficients:
## (Intercept)      x1$nox
##        41.35     -33.92

```

```

model3=lm(y~x1$black)
model3

##
## Call:
## lm(formula = y ~ x1$black)
##
## Coefficients:
## (Intercept)      x1$black
##       10.55103     0.03359

model4=lm(y~x1$crim)
model4

##
## Call:
## lm(formula = y ~ x1$crim)
##
## Coefficients:
## (Intercept)      x1$crim
##       24.0331     -0.4152

models=c("model of y vs crim","model of y vs nox","model of y vs
black","model of y vs lstat")
beta_hat_0=c(coef(model1)[1],coef(model2)[1],coef(model3)[1],coef(model4)[1])
beta_hat=c(coef(model1)[2],coef(model2)[2],coef(model3)[2],coef(model4)[2])
Data=data.frame(models,beta_hat_0,beta_hat)
Data

##           models beta_hat_0    beta_hat
## 1 model of y vs crim  24.03311 -0.41519028
## 2 model of y vs nox   41.34587 -33.91605501
## 3 model of y vs black 10.55103   0.03359306
## 4 model of y vs lstat 24.03311 -0.41519028

summary(model1)

##
## Call:
## lm(formula = y ~ x1$crim)
##
## Residuals:
##      Min      1Q      Median      3Q      Max 
## -16.957  -5.449   -2.007    2.512   29.800 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 24.03311   0.40914   58.74 <2e-16 ***
## x1$crim     -0.41519   0.04389   -9.46 <2e-16 ***
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

## 
## Residual standard error: 8.484 on 504 degrees of freedom
## Multiple R-squared:  0.1508, Adjusted R-squared:  0.1491
## F-statistic: 89.49 on 1 and 504 DF,  p-value: < 2.2e-16

summary(model2)

## 
## Call:
## lm(formula = y ~ x1$nox)
## 
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -13.691  -5.121  -2.161   2.959  31.310 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept)  41.346     1.811   22.83 <2e-16 ***
## x1$nox      -33.916     3.196  -10.61 <2e-16 ***
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 

## 
## Residual standard error: 8.323 on 504 degrees of freedom
## Multiple R-squared:  0.1826, Adjusted R-squared:  0.181 
## F-statistic: 112.6 on 1 and 504 DF,  p-value: < 2.2e-16

summary(model3)

## 
## Call:
## lm(formula = y ~ x1$black)
## 
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -18.884  -4.862  -1.684   2.932  27.763 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 10.551034   1.557463   6.775 3.49e-11 ***
## x1$black    0.033593   0.004231   7.941 1.32e-14 *** 
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 

## 
## Residual standard error: 8.679 on 504 degrees of freedom
## Multiple R-squared:  0.1112, Adjusted R-squared:  0.1094 
## F-statistic: 63.05 on 1 and 504 DF,  p-value: 1.318e-14

summary(model4)

## 
## Call:

```

```

## lm(formula = y ~ x1$crim)
##
## Residuals:
##   Min     1Q Median     3Q    Max
## -16.957 -5.449 -2.007  2.512 29.800
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 24.03311   0.40914  58.74 <2e-16 ***
## x1$crim     -0.41519   0.04389  -9.46 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.484 on 504 degrees of freedom
## Multiple R-squared:  0.1508, Adjusted R-squared:  0.1491
## F-statistic: 89.49 on 1 and 504 DF,  p-value: < 2.2e-16

```

None of them are a good fit but model 1 works better compared to the other models since the R square values are very less

Comparing predictors from different models

```

summary(model1)

##
## Call:
## lm(formula = y ~ x1$crim)
##
## Residuals:
##   Min     1Q Median     3Q    Max
## -16.957 -5.449 -2.007  2.512 29.800
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 24.03311   0.40914  58.74 <2e-16 ***
## x1$crim     -0.41519   0.04389  -9.46 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.484 on 504 degrees of freedom
## Multiple R-squared:  0.1508, Adjusted R-squared:  0.1491
## F-statistic: 89.49 on 1 and 504 DF,  p-value: < 2.2e-16

```

```
summary(model2)
```

```

##
## Call:
## lm(formula = y ~ x1$nox)
##
## Residuals:
##   Min     1Q Median     3Q    Max
## -16.957 -5.449 -2.007  2.512 29.800
## 
```

```

## -13.691 -5.121 -2.161  2.959 31.310
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 41.346     1.811   22.83 <2e-16 ***
## x1$nox     -33.916     3.196  -10.61 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.323 on 504 degrees of freedom
## Multiple R-squared:  0.1826, Adjusted R-squared:  0.181
## F-statistic: 112.6 on 1 and 504 DF,  p-value: < 2.2e-16

summary(model3)

##
## Call:
## lm(formula = y ~ x1$black)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -18.884 -4.862 -1.684  2.932 27.763
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 10.551034  1.557463  6.775 3.49e-11 ***
## x1$black     0.033593  0.004231  7.941 1.32e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.679 on 504 degrees of freedom
## Multiple R-squared:  0.1112, Adjusted R-squared:  0.1094
## F-statistic: 63.05 on 1 and 504 DF,  p-value: 1.318e-14

summary(model4)

##
## Call:
## lm(formula = y ~ x1$crim)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -16.957 -5.449 -2.007  2.512 29.800
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 24.03311  0.40914  58.74 <2e-16 ***
## x1$crim     -0.41519  0.04389  -9.46 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##

```

```
## Residual standard error: 8.484 on 504 degrees of freedom
## Multiple R-squared:  0.1508, Adjusted R-squared:  0.1491
## F-statistic: 89.49 on 1 and 504 DF,  p-value: < 2.2e-16
```